

# Half-precision floating-point format

In [computing](#), **half precision** is a [binary floating-point computer number format](#) that occupies [16 bits](#) (two bytes in modern computers) in [computer memory](#).

In the [IEEE 754-2008](#) standard, the 16-bit [base-2](#) format is referred to as **binary16**. It is intended for storage of floating-point values in applications where higher precision is not essential for performing arithmetic computations.

Although implementations of the IEEE Half-precision floating point are relatively new, several earlier 16-bit floating point formats have existed including that of Hitachi's HD61810 DSP<sup>[1]</sup> of 1982, Scott's WIF<sup>[2]</sup> and the [3dfx Voodoo Graphics processor](#).<sup>[3]</sup>

[Nvidia](#) and [Microsoft](#) defined the **half** datatype in the [Cg language](#), released in early 2002, and implemented it in silicon in the [GeForce FX](#), released in late 2002.<sup>[4]</sup> [ILM](#) was searching for an image format that could handle a wide [dynamic range](#), but without the hard drive and memory cost of floating-point representations that are commonly used for floating-point computation (single and double precision).<sup>[5]</sup> The hardware-accelerated programmable shading group led by John Airey at [SGI \(Silicon Graphics\)](#) invented the s10e5 data type in 1997 as part of the 'bali' design effort. This is described in a [SIGGRAPH 2000 paper](#)<sup>[6]</sup> (see section 4.3) and further documented in US patent 7518615.<sup>[7]</sup>

This format is used in several [computer graphics](#) environments including [MATLAB](#), [OpenEXR](#), [JPEG XR](#), [GIMP](#), [OpenGL](#), [Cg](#), [Direct3D](#), and [D3DX](#). The advantage over 8-bit or 16-bit binary integers is that the increased [dynamic range](#) allows for more detail to be preserved in highlights and [shadows](#) for images. The advantage over 32-bit [single-precision](#) binary formats is that it requires half the storage and [bandwidth](#) (at the expense of precision and range).<sup>[5]</sup>

The [F16C](#) extension allows x86 processors to convert half-precision floats to and from [single-precision floats](#).

Depending on the computer half-precision can be over an order of magnitude faster than double precision, e.g. 37 PFLOPS vs. for half 550 "AI-PFLOPS (Half Precision)".<sup>[8]</sup>

## Contents

### IEEE 754 half-precision binary floating-point format: binary16

- Exponent encoding
- Half precision examples
- Precision limitations on decimal values in [0, 1]
- Precision limitations on decimal values in [1, 2048]
- Precision limitations on integer values

### ARM alternative half-precision

### See also

### References

## Further reading


## External links

# IEEE 754 half-precision binary floating-point format: binary16

The IEEE 754 standard specifies a **binary16** as having the following format:

- Sign bit: 1 bit
- Exponent width: 5 bits
- Significand precision: 11 bits (10 explicitly stored)

The format is laid out as follows:



The format is assumed to have an implicit lead bit with value 1 unless the exponent field is stored with all zeros. Thus only 10 bits of the significand appear in the memory format but the total precision is 11 bits. In IEEE 754 parlance, there are 10 bits of significand, but there are 11 bits of significand precision ( $\log_{10}(2^{11}) \approx 3.311$  decimal digits, or 4 digits  $\pm$  slightly less than 5 units in the last place).

## Exponent encoding

The half-precision binary floating-point exponent is encoded using an offset-binary representation, with the zero offset being 15; also known as exponent bias in the IEEE 754 standard.

- $E_{\min} = 00001_2 - 01111_2 = -14$
- $E_{\max} = 11110_2 - 01111_2 = 15$
- Exponent bias =  $01111_2 = 15$

Thus, as defined by the offset binary representation, in order to get the true exponent the offset of 15 has to be subtracted from the stored exponent.

The stored exponents  $00000_2$  and  $11111_2$  are interpreted specially.

Exponent	Significand = zero	Significand $\neq$ zero	Equation
$00000_2$	<u>zero</u> , <u>-0</u>	<u>subnormal numbers</u>	$(-1)^{\text{signbit}} \times 2^{-14} \times 0.\text{significantbits}_2$
$00001_2, \dots, 11110_2$	normalized value		$(-1)^{\text{signbit}} \times 2^{\text{exponent}-15} \times 1.\text{significantbits}_2$
$11111_2$	<u><math>\pm</math>infinity</u>	<u>NaN</u> (quiet, signalling)	

The minimum strictly positive (subnormal) value is  $2^{-24} \approx 5.96 \times 10^{-8}$ . The minimum positive normal value is  $2^{-14} \approx 6.10 \times 10^{-5}$ . The maximum representable value is  $(2-2^{-10}) \times 2^{15} = 65504$ .

## Half precision examples

These examples are given in bit representation of the floating-point value. This includes the sign bit, (biased) exponent, and significand.

$$0\ 00000\ 00000000001_2 = 0001_{16} = 2^{-14} \times (0 + \frac{1}{1024}) \approx 0.000000059605$$

(smallest positive subnormal number)

$$0\ 00000\ 1111111111_2 = 03ff_{16} = 2^{-14} \times (0 + \frac{1023}{1024}) \approx 0.000060976$$

(largest subnormal number)

$$0\ 00001\ 0000000000_2 = 0400_{16} = 2^{-14} \times (1 + \frac{0}{1024}) \approx 0.000061035$$

(smallest positive normal number)

$$0\ 11110\ 1111111111_2 = 7bff_{16} = 2^{15} \times (1 + \frac{1023}{1024}) = 65504$$

(largest normal number)

$$0\ 01110\ 1111111111_2 = 3bff_{16} = 2^{-1} \times (1 + \frac{1023}{1024}) \approx 0.99951$$

(largest number less than one)

$$0\ 01111\ 0000000000_2 = 3c00_{16} = 2^0 \times (1 + \frac{0}{1024}) = 1$$

(one)

$$0\ 01111\ 0000000001_2 = 3c01_{16} = 2^0 \times (1 + \frac{1}{1024}) \approx 1.001$$

(smallest number larger than one)

$$0\ 01101\ 0101010101_2 = 3555_{16} = 2^{-2} \times (1 + \frac{341}{1024}) = 0.333251953125$$

(equal to 1/3)

$$1\ 10000\ 0000000000_2 = c000_{16} = -2$$

$$0\ 00000\ 0000000000_2 = 0000_{16} = 0$$

$$1\ 00000\ 0000000000_2 = 8000_{16} = -0$$

$$0\ 11111\ 0000000000_2 = 7c00_{16} = \text{infinity}$$

$$1\ 11111\ 0000000000_2 = fc00_{16} = -\text{infinity}$$

By default, 1/3 rounds down like for double precision, because of the odd number of bits in the significand. So the bits beyond the rounding point are 0101... which is less than 1/2 of a unit in the last place.

## Precision limitations on decimal values in [0, 1]

- Decimals between  $2^{-24}$  (minimum positive subnormal) and  $2^{-14}$  (maximum subnormal): fixed interval  $2^{-24}$
- Decimals between  $2^{-14}$  (minimum positive normal) and  $2^{-13}$ : fixed interval  $2^{-24}$

- Decimals between  $2^{-13}$  and  $2^{-12}$ : fixed interval  $2^{-23}$
- Decimals between  $2^{-12}$  and  $2^{-11}$ : fixed interval  $2^{-22}$
- Decimals between  $2^{-11}$  and  $2^{-10}$ : fixed interval  $2^{-21}$
- Decimals between  $2^{-10}$  and  $2^{-9}$ : fixed interval  $2^{-20}$
- Decimals between  $2^{-9}$  and  $2^{-8}$ : fixed interval  $2^{-19}$
- Decimals between  $2^{-8}$  and  $2^{-7}$ : fixed interval  $2^{-18}$
- Decimals between  $2^{-7}$  and  $2^{-6}$ : fixed interval  $2^{-17}$
- Decimals between  $2^{-6}$  and  $2^{-5}$ : fixed interval  $2^{-16}$
- Decimals between  $2^{-5}$  and  $2^{-4}$ : fixed interval  $2^{-15}$
- Decimals between  $2^{-4}$  and  $2^{-3}$ : fixed interval  $2^{-14}$
- Decimals between  $2^{-3}$  and  $2^{-2}$ : fixed interval  $2^{-13}$
- Decimals between  $2^{-2}$  and  $2^{-1}$ : fixed interval  $2^{-12}$
- Decimals between  $2^{-1}$  and  $2^0$ : fixed interval  $2^{-11}$

## Precision limitations on decimal values in [1, 2048]

- Decimals between 1 and 2: fixed interval  $2^{-10}$  ( $1+2^{-10}$  is the next largest float after 1)
- Decimals between 2 and 4: fixed interval  $2^{-9}$
- Decimals between 4 and 8: fixed interval  $2^{-8}$
- Decimals between 8 and 16: fixed interval  $2^{-7}$
- Decimals between 16 and 32: fixed interval  $2^{-6}$
- Decimals between 32 and 64: fixed interval  $2^{-5}$
- Decimals between 64 and 128: fixed interval  $2^{-4}$
- Decimals between 128 and 256: fixed interval  $2^{-3}$
- Decimals between 256 and 512: fixed interval  $2^{-2}$
- Decimals between 512 and 1024: fixed interval  $2^{-1}$
- Decimals between 1024 and 2048: fixed interval  $2^0$

## Precision limitations on integer values

- Integers between 0 and 2048 can be exactly represented (and also between -2048 and 0)
- Integers between 2048 and 4096 round to a multiple of 2 (even number)
- Integers between 4096 and 8192 round to a multiple of 4
- Integers between 8192 and 16384 round to a multiple of 8
- Integers between 16384 and 32768 round to a multiple of 16
- Integers between 32768 and 65519 round to a multiple of  $32^{[9]}$
- Integers above 65519 are rounded to "infinity" if using round-to-even, or above 65535 if using round-to-zero, or above 65504 if using round-to-infinity.

## ARM alternative half-precision

---

ARM processors support (via a floating point control register bit) an "alternative half-precision" format, which does away with the special case for an exponent value of 31 (11111<sub>2</sub>).<sup>[10]</sup> It is almost identical to the IEEE format, but there is no encoding for infinity or NaNs; instead, an exponent of 31 encodes normalized numbers in the range 65536 to 131008.

## See also

---

- bfloat16 floating-point format: Alternative 16-bit floating-point format with 8 bits of exponent and 7 bits of mantissa
- IEEE 754: IEEE standard for floating-point arithmetic (IEEE 754)
- ISO/IEC 10967, Language Independent Arithmetic
- Primitive data type
- RGBE image format

## References

---

1. "hitachi :: dataBooks :: HD61810 Digital Signal Processor Users Manual" ([https://archive.org/details/bitsavers\\_hitachidat1SignalProcessorUsersManual\\_4735688](https://archive.org/details/bitsavers_hitachidat1SignalProcessorUsersManual_4735688)). *Archive.org*. Retrieved 2017-07-14.
2. Scott, Thomas J. (March 1991). "Mathematics and Computer Science at Odds over Real Numbers" (<https://dl.acm.org/citation.cfm?id=107029>). *SIGCSE '91 Proceedings of the twenty-second SIGCSE technical symposium on Computer science education*. **23** (1): 130–139.
3. "/home/usr/bk/glide/docs2.3.1/GLIDEPGM.DOC" (<http://www.gamers.org/dEngine/xf3D/glide/glidepgm.htm>). *Gamers.org*. Retrieved 2017-07-14.
4. "vs\_2\_sw" ([https://developer.download.nvidia.com/cg/vs\\_2\\_sw.html](https://developer.download.nvidia.com/cg/vs_2_sw.html)). *Cg 3.1 Toolkit Documentation*. Nvidia. Retrieved 17 August 2016.
5. "OpenEXR" (<http://www.openexr.com/about.html>). OpenEXR. Retrieved 2017-07-14.
6. Mark S. Peercy; Marc Olano; John Airey; P. Jeffrey Ungar. "Interactive Multi-Pass Programmable Shading" (<https://people.csail.mit.edu/ericchan/bib/pdf/p425-peercy.pdf>) (PDF). *People.csail.mit.edu*. Retrieved 2017-07-14.
7. "Patent US7518615 - Display system having floating point rasterization and floating point ... - Google Patents" (<https://www.google.com/patents/US7518615>). *Google.com*. Retrieved 2017-07-14.
8. "About ABCI - About ABCI | ABCI" ([https://abci.ai/en/about\\_abci/](https://abci.ai/en/about_abci/)). *abci.ai*. Retrieved 2019-10-06.
9. "Medium float calculator" (<https://oletus.github.io/float16-simulator.js/>). Retrieved 2016-07-26. Half precision floating point calculator
10. "Half-precision floating-point number support" (<http://infocenter.arm.com/help/topic/com.arm.doc.dui0205j/CIHGAECI.html>). *RealView Compilation Tools Compiler User Guide*. 10 December 2010. Retrieved 2015-05-05.

## Further reading

---

- Khronos Vulkan signed 16-bit floating point format (<https://www.khronos.org/registry/DataFormat/specs/1.2/dataformat.1.2.html#16bitfp>)

## External links

---

- [Minifloats \(https://www.mrob.com/pub/math/floatformats.html#minifloat\)](https://www.mrob.com/pub/math/floatformats.html#minifloat) (in *Survey of Floating-Point Formats*)
- [OpenEXR site \(http://www.openexr.org/\)](http://www.openexr.org/)
- [Half precision constants \(https://technet.microsoft.com/en-us/library/bb147247\(v=vs.85\).aspx\)](https://technet.microsoft.com/en-us/library/bb147247(v=vs.85).aspx) from [D3DX](#)
- [OpenGL treatment of half precision \(https://web.archive.org/web/20170531074746/http://oss.sgi.com/projects/ogl-sample/registry/ARB/half\\_float\\_pixel.txt\)](https://web.archive.org/web/20170531074746/http://oss.sgi.com/projects/ogl-sample/registry/ARB/half_float_pixel.txt)
- [Fast Half Float Conversions \(http://www.fox-toolkit.org/ftp/fasthalffloatconversion.pdf\)](http://www.fox-toolkit.org/ftp/fasthalffloatconversion.pdf)
- [Analog Devices variant \(http://www.analog.com/static/imported-files/processor\\_manuals/ADSP\\_2136x\\_PGR\\_rev1-1.pdf\)](http://www.analog.com/static/imported-files/processor_manuals/ADSP_2136x_PGR_rev1-1.pdf) (four-bit exponent)
- [C source code to convert between IEEE double, single, and half precision can be found here \(https://www.mathworks.com/matlabcentral/fileexchange/23173\)](https://www.mathworks.com/matlabcentral/fileexchange/23173)
- [Java source code for half-precision floating-point conversion \(https://stackoverflow.com/a/6162687/237321\)](https://stackoverflow.com/a/6162687/237321)
- [Half precision floating point for one of the extended GCC features \(https://gcc.gnu.org/onlinedocs/gcc/Half-Precision.html\)](https://gcc.gnu.org/onlinedocs/gcc/Half-Precision.html)

---

Retrieved from "[https://en.wikipedia.org/w/index.php?title=Half-precision\\_floating-point\\_format&oldid=934602723](https://en.wikipedia.org/w/index.php?title=Half-precision_floating-point_format&oldid=934602723)"

---

**This page was last edited on 7 January 2020, at 13:06 (UTC).**

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. By using this site, you agree to the [Terms of Use](#) and [Privacy Policy](#). Wikipedia® is a registered trademark of the [Wikimedia Foundation, Inc.](#), a non-profit organization.